

```
In [ ]: # import libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# import scalers
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import RobustScaler
from sklearn.preprocessing import Normalizer

# import feature selection
from sklearn.feature_selection import VarianceThreshold
from sklearn.feature_selection import chi2
from sklearn.feature_selection import f_classif
from sklearn.feature_selection import mutual_info_classif
from sklearn.feature_selection import mutual_info_regression
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import SelectPercentile
from sklearn.feature_selection import SelectFromModel
from sklearn.feature_selection import RFE

# import classifiers
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier

# import regressors
from sklearn.linear_model import LinearRegression
from sklearn.svm import SVR
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor

# import metrics
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import r2_score
from sklearn.metrics import classification_report
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from sklearn.metrics import confusion_matrix

# import library for unalanced data

# Synthetic Minority Oversampling Technique (SMOTE)
from imblearn.over_sampling import SMOTE
# K-Nearest Neighbor OverSampling (KNNOR)
from knnorr import data_augment
# SMOTE + Tomek
from imblearn.combine import SMOTETomek
# SMOTE + ENN
from imblearn.combine import SMOTEENN
# random over sampler
```

```
from imblearn.over_sampling import RandomOverSampler

# import system
import os
import sys
```

Deal with unbalanced data

<https://www.analyticsvidhya.com/blog/2020/10/overcoming-class-imbalance-using-smote-techniques/>

K-Nearest Neighbor OverSampling approach:

<https://www.sciencedirect.com/science/article/pii/S156849462>

Online learning:

<https://www.sciencedirect.com/topics/physics-and-astronomy/weight-vector>

```
In [ ]: models_classific = {
        'LogisticRegression': LogisticRegression(),
        'SVC': SVC(),
        'DecisionTreeClassifier': DecisionTreeClassifier(),
        'RandomForestClassifier': RandomForestClassifier(),
    }
```

```
In [ ]: def get_data(PATH, file_name):
        try:
            df = pd.read_excel(PATH+file_name, header=1)
            # fill empty space with _
            csv_filename = file_name.replace(" ", "_")
            df.to_csv(PATH+csv_filename[:-5]+".csv", index=False)
            df = pd.read_csv(PATH+csv_filename[:-5]+".csv")
            return df
        except:
            print("Error: file not found")
            sys.exit(1)
```

```
In [ ]: df = get_data("../data/", "default of credit card clients.xls")
df.head()
```

Out []:

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4
0	1	20000	2	2	1	24	2	2	-1	
1	2	120000	2	2	2	26	-1	2	0	
2	3	90000	2	2	2	34	0	0	0	
3	4	50000	2	2	1	37	0	0	0	
4	5	50000	1	2	1	57	-1	0	-1	

5 rows x 25 columns

```
In [ ]: # make function looks nicer
def data_summary(df, interactive=False):
    """
    Prints a summary of the given DataFrame.

    Parameters:
    df (pd.DataFrame): The DataFrame to summarize.
    interactive (bool): If True, pauses after each summary part and clear
                        the console.

    Returns:
    dict: A dictionary containing various summary information of the Data
    """
    hashtable = {
        "Data shape": df.shape,
        "Data columns": df.columns.to_list(),
        "Data types": df.dtypes.to_dict(),
        "Data describe": df.describe().to_string(),
        "Data null count": df.isnull().sum().to_dict(),
        "Data Count": df.count().to_dict()
    }

    for key, value in hashtable.items():
        print(f"{key}:\n{value}\n")
        if interactive:
            input("Press Enter to continue...")
            os.system('cls' if os.name == 'nt' else 'clear')

    return None
```


```
In [ ]: # check if the data is clean enough
def check_data(df):
    return df.isnull().sum()

# check if the the range of each column, not include the first row
def check_range(df):
    for col in df.columns:
        print(col, df[col].unique())
```

```
In [ ]: # check the correlation between each column
corr = df.corr()
plt.figure(figsize=(14, 14))
```

```
plt.title('Correlation Matrix')
sns.heatmap(corr, annot=True, vmin=-1, vmax=1, cmap='coolwarm')
```

Out[]: <AxesSubplot:title={'center':'Correlation Matrix'}>



```
In [ ]: # write a function about one-hot encoding
def one_hot_encoding(df, col_list):
    df = df.copy()
    for col in col_list:
        dummies = pd.get_dummies(df[col], prefix=col[:4])
        df = pd.concat([df, dummies], axis=1)
        df = df.drop(col, axis=1)
    return df
```

```
In [ ]: def pre_processing_class(df, scaler):
    # remove unnecessary columns
    df = df.drop(['ID'], axis=1)
    # One-hot encoding
    col_list = ["EDUCATION", "MARRIAGE"]
    df = one_hot_encoding(df, col_list)
    # scale the data
    X = df.drop(['default payment next month'], axis=1)
    X = pd.DataFrame(scaler.fit_transform(X), columns=X.columns)
```

```
y = df['default payment next month']
return X,y
```

Task1: Build a classification model that predicts whether or not a customer will default on their next payment

```
In [ ]: # scaler name
scaler_standard = StandardScaler()
scaler_min_max = MinMaxScaler()
scaler_robust = RobustScaler()
scaler_norm = Normalizer()
X, y = pre_processing_class(df, scaler_standard)

In [ ]: def train_test_split_class(X,y,random_state=33,test_size=0.2,oversample=F
# split the data
if cross_val:
    pass
    # do something
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=t

if oversample:
    X_train_oversample, y_train_oversample = oversampler.fit_resample
else:
    X_train_oversample, y_train_oversample = [],[]
return X_train, X_test, y_train, y_test,X_train_oversample, y_train_o

In [ ]: def train_models(name,model, X_train, y_train):
model.fit(X_train, y_train)
print(name + ' trained.')
print('Training accuracy: {:.2f}%'.format(accuracy_score(y_train, mod

def evaluate_models(name,model, X_test, y_test):

# accuracy
print(name + ' Accuracy: {:.2f}%'.format(model.score(X_test, y_test))

# ROC AUC
if name != 'SVC':
    print(name + ' ROC AUC: {:.2f}%'.format(roc_auc_score(y_test, mod

# confusion matrix
y_pred = model.predict(X_test)
cm = confusion_matrix(y_test, y_pred)
recall = cm[1][1] / (cm[1][1] + cm[1][0])
precision = cm[1][1] / (cm[1][1] + cm[0][1])
specificity = cm[0][0] / (cm[0][0] + cm[0][1])
print(name + ' Recall: {:.2f}%'.format(recall * 100))
print(name + ' Precision: {:.2f}%'.format(precision * 100))
print(name + ' Specificity: {:.2f}%'.format(specificity * 100))

# confusion matrix report
print(name + ' Confusion Matrix Report: \n', classification_report(y_

if name != 'SVC':
    # ROC curve
    fpr, tpr, _ = roc_curve(y_test, model.predict_proba(X_test)[:,:1])
```

```

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=name)
plt.plot([0, 1], [0, 1], color='black', linestyle='--')
plt.xlabel('False Positive Rate (Fall-Out)')
plt.ylabel('True Positive Rate (Recall)')
plt.title('ROC Curve')
plt.legend()
plt.show()

# test overfitting
def cross_validation(name,model, X, y):
    scores = cross_val_score(model, X, y, cv=5)
    print(name + ' Cross Validation Accuracy: {:.2f}%'.format(scores.mean

```

```

In [ ]: def model_start(model_name,X_train,y_train,X,y,X_test,y_test):
    train_models(model_name,models_classific[model_name], X_train, y_train)
    evaluate_models(model_name,models_classific[model_name], X_test, y_test)
    cross_validation(model_name,models_classific[model_name], X, y)

def train_different_sampler(X,y,modelname,random_state=33):

    sampler = {
        'SMOTE': SMOTE(random_state=random_state),
        'SMOTETomek': SMOTETomek(random_state=random_state),
        'SMOTEENN': SMOTEENN(random_state=random_state),
        'RandomOverSampler': RandomOverSampler(random_state=random_state)
        # 'Knnor': data_augment.KNNOR()
    }

    print("-----")

    print("Training without oversampling")
    X_train, X_test, y_train, y_test, X_train_oversample, y_train_oversample = \
    model_start(modelname,X_train,y_train,X,y,X_test,y_test)

    print("-----")

    print("Oversampler Using Different Sampler")

    print("-----")

    for name, sampler in sampler.items():
        print("OverSampling with " + name)
        X_train, X_test, y_train, y_test, X_train_oversample, y_train_oversample = \
        # model_name,X_train,y_train,X,y,X_test,y_test
        model_start(modelname,X_train_oversample,y_train_oversample,X,y,X_test,y_test)
        print("-----")

```

Logistic Regression

```


In [ ]: train_different_sampler(X,y,'LogisticRegression')

```

```
-----
Training without oversampling
LogisticRegression trained.
Training accuracy: 81.11%
LogisticRegression Accuracy: 81.22%
LogisticRegression ROC AUC: 72.42%
LogisticRegression Recall: 23.68%
LogisticRegression Precision: 72.62%
LogisticRegression Specificity: 97.48%
LogisticRegression Confusion Matrix Report:
              precision    recall  f1-score   support

      0       0.82        0.97        0.89       4678
      1       0.73        0.24        0.36       1322

 accuracy          0.81        6000
 macro avg         0.77        0.61        0.62        6000
 weighted avg      0.80        0.81        0.77        6000
```




LogisticRegression Cross Validation Accuracy: 81.03%

Oversampler Using Different Sampler

OverSampling with SMOTE
LogisticRegression trained.
Training accuracy: 67.59%
LogisticRegression Accuracy: 66.88%
LogisticRegression ROC AUC: 72.62%
LogisticRegression Recall: 65.36%
LogisticRegression Precision: 36.11%
LogisticRegression Specificity: 67.32%
LogisticRegression Confusion Matrix Report:


	precision	recall	f1-score	support
0	0.87	0.67	0.76	4678
1	0.36	0.65	0.47	1322
accuracy			0.67	6000
macro avg	0.62	0.66	0.61	6000
weighted avg	0.76	0.67	0.70	6000



LogisticRegression Cross Validation Accuracy: 81.03%

OverSampling with SMOTETomek
LogisticRegression trained.
Training accuracy: 68.05%
LogisticRegression Accuracy: 67.00%
LogisticRegression ROC AUC: 72.63%
LogisticRegression Recall: 65.28%
LogisticRegression Precision: 36.20%
LogisticRegression Specificity: 67.49%
LogisticRegression Confusion Matrix Report:

	precision	recall	f1-score	support
0	0.87	0.67	0.76	4678
1	0.36	0.65	0.47	1322
accuracy			0.67	6000
macro avg	0.62	0.66	0.61	6000
weighted avg	0.76	0.67	0.70	6000



LogisticRegression Cross Validation Accuracy: 81.03%

OverSampling with SMOTEENN

LogisticRegression trained.

Training accuracy: 71.31%

LogisticRegression Accuracy: 55.00%

LogisticRegression ROC AUC: 72.63%


LogisticRegression Recall: 75.95%

LogisticRegression Precision: 29.65%

LogisticRegression Specificity: 49.08%

LogisticRegression Confusion Matrix Report:

	precision	recall	f1-score	support
0	0.88	0.49	0.63	4678
1	0.30	0.76	0.43	1322
accuracy			0.55	6000
macro avg	0.59	0.63	0.53	6000
weighted avg	0.75	0.55	0.58	6000




LogisticRegression Cross Validation Accuracy: 81.03%

OverSampling with RandomOverSampler
LogisticRegression trained.

Training accuracy: 67.17%
LogisticRegression Accuracy: 67.53%
LogisticRegression ROC AUC: 72.65%
LogisticRegression Recall: 65.81%
LogisticRegression Precision: 36.77%
LogisticRegression Specificity: 68.02%

LogisticRegression Confusion Matrix Report:

		precision	recall	f1-score	support
	0	0.88	0.68	0.77	4678
	1	0.37	0.66	0.47	1322
	accuracy			0.68	6000
	macro avg	0.62	0.67	0.62	6000
	weighted avg	0.76	0.68	0.70	6000



LogisticRegression Cross Validation Accuracy: 81.03%

SVM

```
In [ ]: train_different_sampler(X,y,'SVC')
```

 Training without oversampling

SVC trained.

Training accuracy: 82.43%

SVC Accuracy: 82.27%

SVC Recall: 32.83%

SVC Precision: 71.15%

SVC Specificity: 96.24%

SVC Confusion Matrix Report:

	precision	recall	f1-score	support
0	0.84	0.96	0.89	4678
1	0.71	0.33	0.45	1322
accuracy			0.82	6000
macro avg	0.77	0.65	0.67	6000
weighted avg	0.81	0.82	0.80	6000

SVC Cross Validation Accuracy: 81.95%

 Oversampler Using Different Sampler

 OverSampling with SMOTE

SVC trained.

Training accuracy: 72.77%

SVC Accuracy: 77.45%

SVC Recall: 57.41%

SVC Precision: 49.00%

SVC Specificity: 83.11%

SVC Confusion Matrix Report:

	precision	recall	f1-score	support
0	0.87	0.83	0.85	4678
1	0.49	0.57	0.53	1322
accuracy			0.77	6000
macro avg	0.68	0.70	0.69	6000
weighted avg	0.79	0.77	0.78	6000

SVC Cross Validation Accuracy: 81.95%

 OverSampling with SMOTETomek

SVC trained.

Training accuracy: 73.36%

SVC Accuracy: 77.40%

SVC Recall: 57.49%

SVC Precision: 48.91%

SVC Specificity: 83.03%

SVC Confusion Matrix Report:

	precision	recall	f1-score	support
0	0.87	0.83	0.85	4678
1	0.49	0.57	0.53	1322
accuracy			0.77	6000
macro avg	0.68	0.70	0.69	6000
weighted avg	0.79	0.77	0.78	6000

SVC Cross Validation Accuracy: 81.95%

OverSampling with SMOTEENN

SVC trained.

Training accuracy: 83.50%

SVC Accuracy: 66.22%

SVC Recall: 72.09%

SVC Precision: 36.50%

SVC Specificity: 64.56%

SVC Confusion Matrix Report:

	precision	recall	f1-score	support
0	0.89	0.65	0.75	4678
1	0.36	0.72	0.48	1322
accuracy			0.66	6000
macro avg	0.63	0.68	0.62	6000
weighted avg	0.78	0.66	0.69	6000

SVC Cross Validation Accuracy: 81.95%

OverSampling with RandomOverSampler

SVC trained.

Training accuracy: 72.63%

SVC Accuracy: 77.28%

SVC Recall: 58.47%

SVC Precision: 48.71%

SVC Specificity: 82.60%

SVC Confusion Matrix Report:

	precision	recall	f1-score	support
0	0.88	0.83	0.85	4678
1	0.49	0.58	0.53	1322
accuracy			0.77	6000
macro avg	0.68	0.71	0.69	6000
weighted avg	0.79	0.77	0.78	6000

SVC Cross Validation Accuracy: 81.95%


Random Forest model

```
In [ ]: train_different_sampler(X,y,'DecisionTreeClassifier')
```

```
-----
Training without oversampling
DecisionTreeClassifier trained.
Training accuracy: 99.93%
DecisionTreeClassifier Accuracy: 72.25%
DecisionTreeClassifier ROC AUC: 61.27%
DecisionTreeClassifier Recall: 41.75%
DecisionTreeClassifier Precision: 38.15%
DecisionTreeClassifier Specificity: 80.87%
DecisionTreeClassifier Confusion Matrix Report:
              precision    recall  f1-score   support

      0       0.83        0.81        0.82        4678
      1       0.38        0.42        0.40        1322

 accuracy          0.72        6000
 macro avg         0.61        0.61        0.61        6000
 weighted avg      0.73        0.72        0.73        6000
```



DecisionTreeClassifier Cross Validation Accuracy: 72.19%

Oversampler Using Different Sampler

OverSampling with SMOTE

DecisionTreeClassifier trained.

Training accuracy: 99.96%

DecisionTreeClassifier Accuracy: 68.93%

DecisionTreeClassifier ROC AUC: 60.68%


DecisionTreeClassifier Recall: 45.99%

DecisionTreeClassifier Precision: 34.58%

DecisionTreeClassifier Specificity: 75.42%

DecisionTreeClassifier Confusion Matrix Report:


	precision	recall	f1-score	support
0	0.83	0.75	0.79	4678
1	0.35	0.46	0.39	1322
accuracy			0.69	6000
macro avg	0.59	0.61	0.59	6000
weighted avg	0.72	0.69	0.70	6000



DecisionTreeClassifier Cross Validation Accuracy: 72.31%

OverSampling with SMOTETomek
DecisionTreeClassifier trained.
Training accuracy: 99.96%
DecisionTreeClassifier Accuracy: 69.15%
DecisionTreeClassifier ROC AUC: 61.16%
DecisionTreeClassifier Recall: 46.90%
DecisionTreeClassifier Precision: 35.05%
DecisionTreeClassifier Specificity: 75.44%
DecisionTreeClassifier Confusion Matrix Report:


	precision	recall	f1-score	support
0	0.83	0.75	0.79	4678
1	0.35	0.47	0.40	1322
accuracy			0.69	6000
macro avg	0.59	0.61	0.60	6000
weighted avg	0.73	0.69	0.71	6000



DecisionTreeClassifier Cross Validation Accuracy: 72.29%

OverSampling with SMOTEENN
DecisionTreeClassifier trained.
Training accuracy: 100.00%
DecisionTreeClassifier Accuracy: 68.13%
DecisionTreeClassifier ROC AUC: 66.02%
DecisionTreeClassifier Recall: 62.25%
DecisionTreeClassifier Precision: 36.81%
DecisionTreeClassifier Specificity: 69.79%
DecisionTreeClassifier Confusion Matrix Report:


	precision	recall	f1-score	support
0	0.87	0.70	0.77	4678
1	0.37	0.62	0.46	1322
accuracy			0.68	6000
macro avg	0.62	0.66	0.62	6000
weighted avg	0.76	0.68	0.71	6000



DecisionTreeClassifier Cross Validation Accuracy: 72.43%

OverSampling with RandomOverSampler
DecisionTreeClassifier trained.
Training accuracy: 99.95%
DecisionTreeClassifier Accuracy: 73.05%
DecisionTreeClassifier ROC AUC: 61.09%
DecisionTreeClassifier Recall: 39.71%
DecisionTreeClassifier Precision: 39.03%
DecisionTreeClassifier Specificity: 82.47%
DecisionTreeClassifier Confusion Matrix Report:

		precision	recall	f1-score	support
	0	0.83	0.82	0.83	4678
	1	0.39	0.40	0.39	1322
	accuracy			0.73	6000
	macro avg	0.61	0.61	0.61	6000
	weighted avg	0.73	0.73	0.73	6000




DecisionTreeClassifier Cross Validation Accuracy: 72.28%

```
In [ ]: train_different_sampler(X,y,'RandomForestClassifier')
```

```
-----
Training without oversampling
RandomForestClassifier trained.
Training accuracy: 99.93%
RandomForestClassifier Accuracy: 81.60%
RandomForestClassifier ROC AUC: 76.43%
RandomForestClassifier Recall: 36.38%
RandomForestClassifier Precision: 64.65%
RandomForestClassifier Specificity: 94.38%
RandomForestClassifier Confusion Matrix Report:
              precision    recall  f1-score   support

      0       0.84        0.94        0.89       4678
      1       0.65        0.36        0.47       1322

 accuracy          0.82       6000
 macro avg         0.74        0.65        0.68       6000
 weighted avg      0.80        0.82        0.80       6000
```




RandomForestClassifier Cross Validation Accuracy: 81.51%

Oversampler Using Different Sampler

OverSampling with SMOTE
RandomForestClassifier trained.
Training accuracy: 99.96%
RandomForestClassifier Accuracy: 79.77%
RandomForestClassifier ROC AUC: 75.85%
RandomForestClassifier Recall: 47.28%
RandomForestClassifier Precision: 54.73%
RandomForestClassifier Specificity: 88.95%
RandomForestClassifier Confusion Matrix Report:


	precision	recall	f1-score	support
0	0.86	0.89	0.87	4678
1	0.55	0.47	0.51	1322
accuracy			0.80	6000
macro avg	0.70	0.68	0.69	6000
weighted avg	0.79	0.80	0.79	6000



RandomForestClassifier Cross Validation Accuracy: 81.36%

OverSampling with SMOTETomek
RandomForestClassifier trained.
Training accuracy: 99.96%
RandomForestClassifier Accuracy: 79.90%
RandomForestClassifier ROC AUC: 75.38%
RandomForestClassifier Recall: 47.43%
RandomForestClassifier Precision: 55.10%
RandomForestClassifier Specificity: 89.08%
RandomForestClassifier Confusion Matrix Report:


	precision	recall	f1-score	support
0	0.86	0.89	0.87	4678
1	0.55	0.47	0.51	1322
accuracy			0.80	6000
macro avg	0.70	0.68	0.69	6000
weighted avg	0.79	0.80	0.79	6000



RandomForestClassifier Cross Validation Accuracy: 81.50%

OverSampling with SMOTEENN
RandomForestClassifier trained.
Training accuracy: 100.00%
RandomForestClassifier Accuracy: 75.27%
RandomForestClassifier ROC AUC: 76.77%
RandomForestClassifier Recall: 61.57%
RandomForestClassifier Precision: 45.47%
RandomForestClassifier Specificity: 79.14%
RandomForestClassifier Confusion Matrix Report:

	precision	recall	f1-score	support
0	0.88	0.79	0.83	4678
1	0.45	0.62	0.52	1322
accuracy			0.75	6000
macro avg	0.67	0.70	0.68	6000
weighted avg	0.79	0.75	0.76	6000



RandomForestClassifier Cross Validation Accuracy: 81.49%

OverSampling with RandomOverSampler

RandomForestClassifier trained.

Training accuracy: 99.95%

RandomForestClassifier Accuracy: 80.88%

RandomForestClassifier ROC AUC: 76.80%


RandomForestClassifier Recall: 42.97%

RandomForestClassifier Precision: 59.11%

RandomForestClassifier Specificity: 91.60%

RandomForestClassifier Confusion Matrix Report:

	precision	recall	f1-score	support
0	0.85	0.92	0.88	4678
1	0.59	0.43	0.50	1322
accuracy			0.81	6000
macro avg	0.72	0.67	0.69	6000
weighted avg	0.79	0.81	0.80	6000



RandomForestClassifier Cross Validation Accuracy: 81.59%

Build a regression model that predicts a customer's limit balance if they were a new customer to the bank with only X0, X2-5 available to you as data.

```
In [ ]: def pre_processing_regre(df, scaler):
        # remove unnecessary columns
        df = df.drop(['ID'], axis=1)
```

```

# One-hot encoding
col_list = ["EDUCATION", "MARRIAGE"]
df = one_hot_encoding(df, col_list)
# scale the data
X = df.drop(['LIMIT_BAL'], axis=1)
X = pd.DataFrame(scaler.fit_transform(X), columns=X.columns)
y = df['LIMIT_BAL']
return X,y

```

```
X, y = pre_processing_regre(df, scaler_standard)
```

```

In [ ]: models_regre = {
        'LinearRegression': LinearRegression(),
        'SVR': SVR(),
        'DecisionTreeRegressor': DecisionTreeRegressor(),
        'RandomForestRegressor': RandomForestRegressor(),
    }

```

```

In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,

def train_models(name,model, X_train, y_train):
    model.fit(X_train, y_train)
    print(name + ' trained.')

def evaluate_models(name,model, X_test, y_test):
    print(name + ' Accuracy: {:.2f}%'.format(model.score(X_test, y_test)

```

Linear Regression

```

In [ ]: train_models("linear_regression",models_regre["LinearRegression"], X_train, y_train)
        evaluate_models("linear_regression",models_regre["LinearRegression"], X_test, y_test)

```

linear_regression trained.
linear_regression Accuracy: 36.10%

Random Forest

```

In [ ]: train_models("SVR",models_regre["SVR"], X_train, y_train)
        evaluate_models("SVR",models_regre["SVR"], X_test, y_test)

```

SVR trained.
SVR Accuracy: -4.24%

```

In [ ]: train_models("DecisionTreeRegressor",models_regre["DecisionTreeRegressor"], X_train, y_train)
        evaluate_models("DecisionTreeRegressor",models_regre["DecisionTreeRegressor"], X_test, y_test)

```

DecisionTreeRegressor trained.
DecisionTreeRegressor Accuracy: -6.16%

```

In [ ]: train_models("RandomForestRegressor",models_regre["RandomForestRegressor"], X_train, y_train)
        evaluate_models("RandomForestRegressor",models_regre["RandomForestRegressor"], X_test, y_test)

```

RandomForestRegressor trained.
RandomForestRegressor Accuracy: 46.90%

ANN model

In []: